

# Learning Structured Predictors

Xavier Carreras



TALP Research Center  
Universitat Politècnica de Catalunya

# Supervised (Structured) Prediction

- ▶ Learning to predict: given training data

$$\{(\mathbf{x}^{(1)}, \mathbf{y}^{(1)}), (\mathbf{x}^{(2)}, \mathbf{y}^{(2)}), \dots, (\mathbf{x}^{(m)}, \mathbf{y}^{(m)})\}$$

learn a predictor  $\mathbf{x} \rightarrow \mathbf{y}$  that *works well* on unseen inputs  $\mathbf{x}$

- ▶ Non-Structured Prediction: outputs  $\mathbf{y}$  are atomic
  - ▶ Binary prediction:  $\mathbf{y} \in \{-1, +1\}$
  - ▶ Multiclass prediction:  $\mathbf{y} \in \{1, 2, \dots, L\}$
- ▶ Structured Prediction: outputs  $\mathbf{y}$  are structured
  - ▶ Sequence prediction:  $\mathbf{y}$  are sequences
  - ▶ Parsing:  $\mathbf{y}$  are trees
  - ▶ ...

# Named Entity Recognition

y	PER	-	QNT	-	-	ORG	ORG	-	TIME
x	Jim	bought	300	shares	of	Acme	Corp.	in	2006

# Named Entity Recognition

y	PER	-	QNT	-	-	ORG	ORG	-	TIME
x	Jim	bought	300	shares	of	Acme	Corp.	in	2006


y	PER	PER	-	-	LOC
x	Jack	London	went	to	Paris

y	PER	PER	-	-	LOC
x	Paris	Hilton	went	to	London

# Part-of-speech Tagging

<b>y</b>	NNP	NNP	VBZ	NNP	.
<b>x</b>	Ms.	Haag	plays	Elianti	.


# Syntactic Parsing



x are sentences

y are syntactic dependency trees

# Machine Translation



(Galley et al 2006)

x are sentences in Chinese

y are sentences in English aligned to x

# Object Detection



(Kumar and Hebert 2003)

$x$  are images  
 $y$  are grids labeled with object types



# Object Detection



(Kumar and Hebert 2003)

$x$  are images  
 $y$  are grids labeled with object types

# Today's Goals

- ▶ Introduce basic tools for structure prediction
  - ▶ We will restrict to sequence prediction
- ▶ Understand what tools we can use from standard classification
  - ▶ Learning paradigms and algorithms, in essence, work here too
  - ▶ However, computations behind algorithms are prohibitive
- ▶ Understand what tools can we use from existing formalisms that model structured data
  - ▶ We will borrow inference algorithms for tractable computations
  - ▶ E.g., algorithms for HMMs (Viterbi, forward-backward) will play a major role in today's methods

# Today's Goals

- ▶ Introduce basic tools for structure prediction
  - ▶ We will restrict to sequence prediction
- ▶ Understand what tools we can use from standard classification
  - ▶ Learning paradigms and algorithms, in essence, work here too
  - ▶ However, computations behind algorithms are prohibitive
- ▶ Understand what tools can we use from existing formalisms that model structured data
  - ▶ We will borrow inference algorithms for tractable computations
  - ▶ E.g., algorithms for HMMs (Viterbi, forward-backward) will play a major role in today's methods

## Sequence Prediction

<b>y</b>	PER	PER	-	-	LOC
<b>x</b>	Jack	London	went	to	Paris

# Sequence Prediction

- ▶  $\mathbf{x} = x_1x_2 \dots x_n$  are input sequences,  $x_i \in \mathcal{X}$
- ▶  $\mathbf{y} = y_1y_2 \dots y_n$  are output sequences,  $y_i \in \{1, \dots, L\}$
- ▶ **Goal:** given training data

$$\{(\mathbf{x}^{(1)}, \mathbf{y}^{(1)}), (\mathbf{x}^{(2)}, \mathbf{y}^{(2)}), \dots, (\mathbf{x}^{(m)}, \mathbf{y}^{(m)})\}$$

learn a predictor  $\mathbf{x} \rightarrow \mathbf{y}$  that **works well** on unseen inputs  $\mathbf{x}$

- ▶ What is the form of our prediction model?

# Approach 1: Local Classifiers

Jack    ?    went    to    Paris

Jack    London    went    to    Paris

Decompose the sequence into  $n$  classification problems:

- ▶ A classifier predicts individual labels at each position

$$\hat{y}_i = \operatorname{argmax}_{l \in \{\text{LOC}, \text{PER}, -\}} \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, i, l)$$

- ▶  $\mathbf{f}(\mathbf{x}, i, l)$  represents an assignment of label  $l$  for  $x_i$
- ▶  $\mathbf{w}$  is a vector of parameters, has a weight for each feature of  $\mathbf{f}$ 
  - ▶ Use standard classification methods to learn  $\mathbf{w}$
- ▶ At test time, predict the best sequence by  
a simple concatenation of the best label for each position

# Approach 1: Local Classifiers

Jack ? went to Paris

Decompose the sequence into  $n$  classification problems:

- ▶ A classifier predicts individual labels at each position

$$\hat{y}_i = \operatorname{argmax}_{l \in \{\text{LOC}, \text{PER}, -\}} \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, i, l)$$

- ▶  $\mathbf{f}(\mathbf{x}, i, l)$  represents an assignment of label  $l$  for  $x_i$
- ▶  $\mathbf{w}$  is a vector of parameters, has a weight for each feature of  $\mathbf{f}$ 
  - ▶ Use standard classification methods to learn  $\mathbf{w}$
- ▶ At test time, predict the best sequence by  
a simple concatenation of the best label for each position

# Indicator Features

- ▶  $\mathbf{f}(\mathbf{x}, i, l)$  is a vector of  $d$  features representing label  $l$  for  $x_i$

$$( \mathbf{f}_1(\mathbf{x}, i, l), \dots, \mathbf{f}_j(\mathbf{x}, i, l), \dots, \mathbf{f}_d(\mathbf{x}, i, l) )$$

- ▶ What's in a feature  $\mathbf{f}_j(\mathbf{x}, i, l)$ ?
  - ▶ Anything we can compute using  $\mathbf{x}$  and  $i$  and  $l$
  - ▶ Anything that indicates whether  $l$  is (not) a good label for  $x_i$
  - ▶ **Indicator features**: binary-valued features looking at a single simple property

$$\mathbf{f}_j(\mathbf{x}, i, l) = \begin{cases} 1 & \text{if } x_i = \text{London and } l = \text{LOC} \\ 0 & \text{otherwise} \end{cases}$$

$$\mathbf{f}_k(\mathbf{x}, i, l) = \begin{cases} 1 & \text{if } x_{i+1} = \text{went and } l = \text{LOC} \\ 0 & \text{otherwise} \end{cases}$$



# More Features for NE Recognition

Jack      <sup>PER</sup>  
London    went    to    Paris

In practice, construct  $\mathbf{f}(\mathbf{x}, i, l)$  by ...

- ▶ Define a number of simple patterns of  $\mathbf{x}$  and  $i$ 
  - ▶ current word  $x_i$
  - ▶ is  $x_i$  capitalized?
  - ▶  $x_i$  has digits?
  - ▶ prefixes/suffixes of size 1, 2, 3, ...
  - ▶ is  $x_i$  a known location?
  - ▶ is  $x_i$  a known person?
  - ▶ next word
  - ▶ previous word
  - ▶ current and next words together
  - ▶ other combinations
- ▶ Generate features by combining patterns with label identities  $l$

# More Features for NE Recognition


PER      PER      -  
Jack   London   went   to   Paris

In practice, construct  $\mathbf{f}(\mathbf{x}, i, l)$  by ...

- ▶ Define a number of simple patterns of  $\mathbf{x}$  and  $i$ 
  - ▶ current word  $x_i$
  - ▶ is  $x_i$  capitalized?
  - ▶  $x_i$  has digits?
  - ▶ prefixes/suffixes of size 1, 2, 3, ...
  - ▶ is  $x_i$  a known location?
  - ▶ is  $x_i$  a known person?
  - ▶ next word
  - ▶ previous word
  - ▶ current and next words together
  - ▶ other combinations
- ▶ Generate features by combining patterns with label identities  $l$

**Main limitation:** features can't capture interactions between labels!

## Approach 2: HMM for Sequence Prediction




- ▶ Define an HMM where each label is a state
- ▶ Model parameters:
  - ▶  $\pi_l$  : probability of starting with label  $l$
  - ▶  $T_{l,l'}$ : probability of transitioning from  $l$  to  $l'$
  - ▶  $O_{l,x}$ : probability of generating symbol  $x$  given label  $l$
- ▶ Predictions:

$$p(\mathbf{x}, \mathbf{y}) = \pi_{y_1} O_{y_1, x_1} \prod_{i>1} T_{y_{i-1}, y_i} O_{y_i, x_i}$$


- ▶ Learning: relative counts + smoothing
- ▶ Prediction: Viterbi algorithm

## Approach 2: Representation in HMM



- ▶ Label interactions are captured in the transition parameters
- ▶ But interaction between labels and input symbols are quite limited!
  - ▶ Only  $O_{y_i, x_i} = p(x_i | y_i)$
  - ▶ Not clear how to exploit patterns such as:
    - ▶ Capitalization, digits
    - ▶ Prefixes and suffixes
    - ▶ Next word, previous word
    - ▶ Combinations of these with label transitions
- ▶ Why? HMM independence assumptions:  
given label  $y_i$ , token  $x_i$  is independent of anything else

## Approach 2: Representation in HMM



- ▶ Label interactions are captured in the transition parameters
- ▶ But interaction between labels and input symbols are quite limited!
  - ▶ Only  $O_{y_i, x_i} = p(x_i | y_i)$
  - ▶ Not clear how to exploit patterns such as:
    - ▶ Capitalization, digits
    - ▶ Prefixes and suffixes
    - ▶ Next word, previous word
    - ▶ Combinations of these with label transitions
- ▶ Why? HMM independence assumptions:  
given label  $y_i$ , token  $x_i$  is independent of anything else

# Local Classifiers vs. HMM

## LOCAL CLASSIFIERS

- ▶ Form:

$$\mathbf{w} \cdot \mathbf{f}(\mathbf{x}, i, l)$$

- ▶ Learning: standard classifiers
- ▶ Prediction: independent for each  $x_i$
- ▶ Advantage: feature-rich
- ▶ Drawback: no label interactions

## HMM

- ▶ Form:

$$\pi_{y_1} O_{y_1, x_1} \prod_{i>1} T_{y_{i-1}, y_i} O_{y_i, x_i}$$

- ▶ Learning: relative counts
- ▶ Prediction: Viterbi
- ▶ Advantage: label interactions
- ▶ Drawback: no fine-grained features

## Approach 3: Global Sequence Predictors

<b>y:</b>	PER	PER	-	-	LOC
<b>x:</b>	Jack	London	went	to	Paris

Learn a single classifier from  $\mathbf{x} \rightarrow \mathbf{y}$

$$\text{predict}(\mathbf{x}_{1:n}) = \operatorname{argmax}_{\mathbf{y} \in \mathcal{Y}^n} \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, \mathbf{y})$$

But ...

- ▶ How do we represent entire sequences in  $\mathbf{f}(\mathbf{x}, \mathbf{y})$ ?
- ▶ There are exponentially-many sequences  $\mathbf{y}$  for a given  $\mathbf{x}$ , how do we solve the  $\operatorname{argmax}$  problem?

## Approach 3: Global Sequence Predictors

<b>y:</b>	PER	PER	-	-	LOC
<b>x:</b>	Jack	London	went	to	Paris

Learn a single classifier from  $\mathbf{x} \rightarrow \mathbf{y}$

$$\text{predict}(\mathbf{x}_{1:n}) = \underset{\mathbf{y} \in \mathcal{Y}^n}{\operatorname{argmax}} \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, \mathbf{y})$$

But ...

- ▶ How do we represent entire sequences in  $\mathbf{f}(\mathbf{x}, \mathbf{y})$ ?
- ▶ There are **exponentially-many** sequences  $\mathbf{y}$  for a given  $\mathbf{x}$ , how do we solve the **argmax** problem?



# Factored Representations

<b>y:</b>	PER	PER	-	-	LOC
<b>x:</b>	Jack	London	went	to	Paris

- ▶ How do we represent entire sequences in  $\mathbf{f}(\mathbf{x}, \mathbf{y})$ ?
  - ▶ Look at individual assignments  $y_i$  (standard classification)
  - ▶ Look at **bigrams** of outputs labels  $\langle y_{i-1}, y_i \rangle$
  - ▶ Look at **trigrams** of outputs labels  $\langle y_{i-2}, y_{i-1}, y_i \rangle$
  - ▶ Look at  **$n$ -grams** of outputs labels  $\langle y_{i-n+1}, \dots, y_{i-1}, y_i \rangle$
  - ▶ Look at the full label sequence  $\mathbf{y}$  (intractable)
- ▶ A factored representation will lead to a tractable model

# Factored Representations

<b>y:</b>	PER	PER	-	-	LOC
<b>x:</b>	Jack	London	went	to	Paris

- ▶ How do we represent entire sequences in  $\mathbf{f}(\mathbf{x}, \mathbf{y})$ ?
  - ▶ Look at individual assignments  $y_i$  (standard classification)
  - ▶ Look at **bigrams** of outputs labels  $\langle y_{i-1}, y_i \rangle$
  - ▶ Look at **trigrams** of outputs labels  $\langle y_{i-2}, y_{i-1}, y_i \rangle$
  - ▶ Look at  **$n$ -grams** of outputs labels  $\langle y_{i-n+1}, \dots, y_{i-1}, y_i \rangle$
  - ▶ Look at the full label sequence  $\mathbf{y}$  (intractable)
- ▶ A factored representation will lead to a tractable model

# Factored Representations

<b>y:</b>	PER	PER	-	-	LOC
<b>x:</b>	Jack	London	went	to	Paris

- ▶ How do we represent entire sequences in  $\mathbf{f}(\mathbf{x}, \mathbf{y})$ ?
  - ▶ Look at individual assignments  $y_i$  (standard classification)
  - ▶ Look at **bigrams** of outputs labels  $\langle y_{i-1}, y_i \rangle$
  - ▶ Look at **trigrams** of outputs labels  $\langle y_{i-2}, y_{i-1}, y_i \rangle$
  - ▶ Look at  **$n$ -grams** of outputs labels  $\langle y_{i-n+1}, \dots, y_{i-1}, y_i \rangle$
  - ▶ Look at the full label sequence  $\mathbf{y}$  (intractable)
- ▶ A factored representation will lead to a tractable model

# Factored Representations

<b>y:</b>	PER	PER	-	-	LOC
<b>x:</b>	Jack	London	went	to	Paris

- ▶ How do we represent entire sequences in  $\mathbf{f}(\mathbf{x}, \mathbf{y})$ ?
  - ▶ Look at individual assignments  $y_i$  (standard classification)
  - ▶ Look at **bigrams** of outputs labels  $\langle y_{i-1}, y_i \rangle$
  - ▶ Look at **trigrams** of outputs labels  $\langle y_{i-2}, y_{i-1}, y_i \rangle$
  - ▶ Look at  **$n$ -grams** of outputs labels  $\langle y_{i-n+1}, \dots, y_{i-1}, y_i \rangle$
  - ▶ Look at the full label sequence  $\mathbf{y}$  (intractable)
- ▶ A factored representation will lead to a tractable model

## Bigram Indicator Features

	1	2	3	4	5
y	PER	PER	-	-	LOC
x	Jack	London	went	to	Paris

- Indicator features:

$$\mathbf{f}_j(\mathbf{x}, i, y_{i-1}, y_i) = \begin{cases} 1 & \text{if } x_i = \text{"London"} \text{ and} \\ & y_{i-1} = \text{PER and } y_i = \text{PER} \\ 0 & \text{otherwise} \end{cases}$$

e.g.,  $\mathbf{f}_j(\mathbf{x}, 2, \text{PER}, \text{PER}) = 1$ ,  $\mathbf{f}_j(\mathbf{x}, 3, \text{PER}, -) = 0$

# More Bigram Indicator Features

	1	2	3	4	5
<b>x</b>	Jack	London	went	to	Paris
<b>y</b>	PER	PER	-	-	LOC
<b>y'</b>	PER	LOC	-	-	LOC
<b>y''</b>	-	-	-	LOC	-
<b>x'</b>	My	trip	to	London	...

$\mathbf{f}_1(\dots) = 1$  iff  $x_i = \text{"London"}$  and  $y_{i-1} = \text{PER}$  and  $y_i = \text{PER}$

$\mathbf{f}_2(\dots) = 1$  iff  $x_i = \text{"London"}$  and  $y_{i-1} = \text{PER}$  and  $y_i = \text{LOC}$

$\mathbf{f}_3(\dots) = 1$  iff  $x_{i-1} \sim /(\text{in}|\text{to}|\text{at})/$  and  $x_i \sim /[\text{A-Z}]/$  and  $y_i = \text{LOC}$

$\mathbf{f}_4(\dots) = 1$  iff  $y_i = \text{LOC}$  and  $\text{WORLD-CITIES}(x_i) = 1$

$\mathbf{f}_5(\dots) = 1$  iff  $y_i = \text{PER}$  and  $\text{FIRST-NAMES}(x_i) = 1$

## More Bigram Indicator Features

	1	2	3	4	5
x	Jack	London	went	to	Paris
y	PER	PER	-	-	LOC
y'	PER	LOC	-	-	LOC
y''	-	-	-	LOC	-
x'	My	trip	to	London	...

$\mathbf{f}_1(\dots) = 1$  iff  $x_i = \text{"London"}$  and  $y_{i-1} = \text{PER}$  and  $y_i = \text{PER}$

$\mathbf{f}_2(\dots) = 1$  iff  $x_i = \text{"London"}$  and  $y_{i-1} = \text{PER}$  and  $y_i = \text{LOC}$

$\mathbf{f}_3(\dots) = 1$  iff  $x_{i-1} \sim /(\text{in}|\text{to}|\text{at})/$  and  $x_i \sim /[\text{A-Z}]/$  and  $y_i = \text{LOC}$

$\mathbf{f}_4(\dots) = 1$  iff  $y_i = \text{LOC}$  and  $\text{WORLD-CITIES}(x_i) = 1$

$\mathbf{f}_5(\dots) = 1$  iff  $y_i = \text{PER}$  and  $\text{FIRST-NAMES}(x_i) = 1$

## More Bigram Indicator Features

	1	2	3	4	5
x	Jack	London	went	to	Paris
y	PER	PER	-	-	LOC
y'	PER	LOC	-	-	LOC
y''	-	-	-	LOC	-
x'	My	trip	to	London	...

$\mathbf{f}_1(\dots) = 1$  iff  $x_i = \text{"London"}$  and  $y_{i-1} = \text{PER}$  and  $y_i = \text{PER}$

$\mathbf{f}_2(\dots) = 1$  iff  $x_i = \text{"London"}$  and  $y_{i-1} = \text{PER}$  and  $y_i = \text{LOC}$

$\mathbf{f}_3(\dots) = 1$  iff  $x_{i-1} \sim /(\text{in}|\text{to}|\text{at})/$  and  $x_i \sim /[\text{A-Z}]/$  and  $y_i = \text{LOC}$

$\mathbf{f}_4(\dots) = 1$  iff  $y_i = \text{LOC}$  and  $\text{WORLD-CITIES}(x_i) = 1$

$\mathbf{f}_5(\dots) = 1$  iff  $y_i = \text{PER}$  and  $\text{FIRST-NAMES}(x_i) = 1$



# More Bigram Indicator Features

	1	2	3	4	5
<b>x</b>	Jack	London	went	to	Paris
<b>y</b>	PER	PER	-	-	LOC
<b>y'</b>	PER	LOC	-	-	LOC
<b>y''</b>	-	-	-	LOC	-
<b>x'</b>	My	trip	to	London	...

$\mathbf{f}_1(\dots) = 1$  iff  $x_i = \text{"London"}$  and  $y_{i-1} = \text{PER}$  and  $y_i = \text{PER}$

$\mathbf{f}_2(\dots) = 1$  iff  $x_i = \text{"London"}$  and  $y_{i-1} = \text{PER}$  and  $y_i = \text{LOC}$

$\mathbf{f}_3(\dots) = 1$  iff  $x_{i-1} \sim /(\text{in}|\text{to}|\text{at})/$  and  $x_i \sim /[\text{A-Z}]/$  and  $y_i = \text{LOC}$

$\mathbf{f}_4(\dots) = 1$  iff  $y_i = \text{LOC}$  and  $\text{WORLD-CITIES}(x_i) = 1$

$\mathbf{f}_5(\dots) = 1$  iff  $y_i = \text{PER}$  and  $\text{FIRST-NAMES}(x_i) = 1$

## More Bigram Indicator Features

	1	2	3	4	5
x	Jack	London	went	to	Paris
y	PER	PER	-	-	LOC
y'	PER	LOC	-	-	LOC
y''	-	-	-	LOC	-
x'	My	trip	to	London	...

$\mathbf{f}_1(\dots) = 1$  iff  $x_i = \text{"London"}$  and  $y_{i-1} = \text{PER}$  and  $y_i = \text{PER}$

$\mathbf{f}_2(\dots) = 1$  iff  $x_i = \text{"London"}$  and  $y_{i-1} = \text{PER}$  and  $y_i = \text{LOC}$

$\mathbf{f}_3(\dots) = 1$  iff  $x_{i-1} \sim /(\text{in}|\text{to}|\text{at})/$  and  $x_i \sim /[\text{A-Z}]/$  and  $y_i = \text{LOC}$

$\mathbf{f}_4(\dots) = 1$  iff  $y_i = \text{LOC}$  and  $\text{WORLD-CITIES}(x_i) = 1$

$\mathbf{f}_5(\dots) = 1$  iff  $y_i = \text{PER}$  and  $\text{FIRST-NAMES}(x_i) = 1$

# Representations Factored at Bigrams

<b>y:</b>	PER	PER	-	-	LOC
<b>x:</b>	Jack	London	went	to	Paris

- ▶  $\mathbf{f}(\mathbf{x}, i, y_{i-1}, y_i)$ 
  - ▶ A  $d$ -dimensional feature vector of a label bigram at  $i$
  - ▶ Each dimension is typically a boolean indicator (0 or 1)
- ▶  $\mathbf{f}(\mathbf{x}, \mathbf{y}) = \sum_{i=1}^n \mathbf{f}(\mathbf{x}, i, y_{i-1}, y_i)$ 
  - ▶ A  $d$ -dimensional feature vector of the entire  $\mathbf{y}$
  - ▶ Aggregated representation by summing bigram feature vectors
  - ▶ Each dimension is now a **count** of a feature pattern

# Linear Sequence Prediction

$$\text{predict}(\mathbf{x}_{1:n}) = \operatorname{argmax}_{\mathbf{y} \in \mathcal{Y}^n} \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, \mathbf{y})$$

where

$$\mathbf{f}(\mathbf{x}, \mathbf{y}) = \sum_{i=1}^n \mathbf{f}(\mathbf{x}, i, y_{i-1}, y_i)$$

- Note the linearity of the expression:

$$\begin{aligned} \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, \mathbf{y}) &= \mathbf{w} \cdot \sum_{i=1}^n \mathbf{f}(\mathbf{x}, i, y_{i-1}, y_i) \\ &= \sum_{i=1}^n \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, i, y_{i-1}, y_i) \end{aligned}$$

- Next questions:
  - How do we solve the **argmax** problem?
  - How do we learn  $\mathbf{w}$ ?

# Linear Sequence Prediction

$$\text{predict}(\mathbf{x}_{1:n}) = \operatorname{argmax}_{\mathbf{y} \in \mathcal{Y}^n} \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, \mathbf{y})$$

where

$$\mathbf{f}(\mathbf{x}, \mathbf{y}) = \sum_{i=1}^n \mathbf{f}(\mathbf{x}, i, y_{i-1}, y_i)$$

- Note the linearity of the expression:

$$\begin{aligned} \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, \mathbf{y}) &= \mathbf{w} \cdot \sum_{i=1}^n \mathbf{f}(\mathbf{x}, i, y_{i-1}, y_i) \\ &= \sum_{i=1}^n \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, i, y_{i-1}, y_i) \end{aligned}$$

- Next questions:
  - How do we solve the **argmax** problem?
  - How do we learn  $\mathbf{w}$ ?

# Linear Sequence Prediction

$$\text{predict}(\mathbf{x}_{1:n}) = \operatorname{argmax}_{\mathbf{y} \in \mathcal{Y}^n} \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, \mathbf{y})$$

where

$$\mathbf{f}(\mathbf{x}, \mathbf{y}) = \sum_{i=1}^n \mathbf{f}(\mathbf{x}, i, y_{i-1}, y_i)$$

- Note the linearity of the expression:

$$\begin{aligned} \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, \mathbf{y}) &= \mathbf{w} \cdot \sum_{i=1}^n \mathbf{f}(\mathbf{x}, i, y_{i-1}, y_i) \\ &= \sum_{i=1}^n \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, i, y_{i-1}, y_i) \end{aligned}$$


- Next questions:
  - How do we solve the [argmax](#) problem?
  - How do we learn  $\mathbf{w}$ ?

# Decoding with CRFs

- Consider a fixed  $\mathbf{w}$ . Given  $\mathbf{x}_{1:n}$  find:

$$\operatorname{argmax}_{\mathbf{y} \in \mathcal{Y}^n} \sum_{i=1}^n \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, i, y_{i-1}, y_i)$$

- We can use the Viterbi algorithm, takes  $O(n|\mathcal{Y}|^2)$
- Intuition: output sequences that share bigrams will share scores



# Viterbi for Linear Factored Predictors

$$\hat{\mathbf{y}} = \operatorname{argmax}_{\mathbf{y} \in \mathcal{Y}^n} \sum_{i=1}^n \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, i, y_{i-1}, y_i)$$

- **Definition:** score of optimal sequence for  $\mathbf{x}_{1:i}$  ending with  $a \in \mathcal{Y}$

$$\delta_i(a) = \max_{\mathbf{y} \in \mathcal{Y}^i: y_i = a} \sum_{j=1}^i \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, j, y_{j-1}, y_j)$$

- Use the following recursions, for all  $a \in \mathcal{Y}$ :

$$\begin{aligned}\delta_1(a) &= \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, 1, y_0 = \text{NULL}, a) \\ \delta_i(a) &= \max_{b \in \mathcal{Y}} \delta_{i-1}(b) + \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, i, b, a)\end{aligned}$$

- The optimal score for  $\mathbf{x}$  is  $\max_{a \in \mathcal{Y}} \delta_n(a)$
- The optimal sequence  $\hat{\mathbf{y}}$  can be recovered through *pointers*



# Linear Factored Sequence Prediction

$$\text{predict}(\mathbf{x}_{1:n}) = \operatorname{argmax}_{\mathbf{y} \in \mathcal{Y}^n} \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, \mathbf{y})$$

- ▶ Factored representation, e.g. based on bigrams
- ▶ Flexible, arbitrary features of full  $\mathbf{x}$  and the factors
- ▶ Efficient prediction using Viterbi
- ▶ Next, learning  $\mathbf{w}$ :
  - ▶ Conditional Random Fields
  - ▶ Perceptron
  - ▶ SVM

# Conditional Random Fields for Sequence Prediction

<b>y</b>	PER	PER	-	-	LOC
<b>x</b>	Jack	London	went	to	Paris

# Conditional Random Fields

(Lafferty, McCallum, Pereira 2001)

- ▶ Model the conditional distribution:

$$P(\mathbf{y}|\mathbf{x}; \mathbf{w})$$

where

- ▶  $\mathbf{x} = x_1 x_2 \dots x_n \in \mathcal{X}^*$
  - ▶  $\mathbf{y} = y_1 y_2 \dots y_n \in \mathcal{Y}^*$  and  $\mathcal{Y} = \{1, \dots, L\}$
  - ▶  $\mathbf{w}$  are model parameters
- 
- ▶ To predict the best sequence

$$\text{predict}(\mathbf{x}_{1:n}) = \operatorname{argmax}_{\mathbf{y} \in \mathcal{Y}^n} P(\mathbf{y}|\mathbf{x})$$

- ▶ We will see it is a particular form of a factored linear predictor

# CRFs are Log-linear models

- ▶ The model form is:

$$P(\mathbf{y}|\mathbf{x}; \mathbf{w}) = \frac{\exp \{ \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, \mathbf{y}) \}}{Z(\mathbf{x}, \mathbf{w})}$$

where

$$Z(\mathbf{x}, \mathbf{w}) = \sum_{\mathbf{z} \in \mathcal{Y}^*} \exp \{ \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, \mathbf{z}) \}$$

- ▶  $\mathbf{f}(\mathbf{x}, \mathbf{y})$  represents  $\mathbf{x}$  and  $\mathbf{y}$  with  $d$  features
- ▶  $\mathbf{w} \in \mathbb{R}^d$  are the parameters of the model
- ▶ CRFs are **log-linear models**:

$$\log P(\mathbf{y}|\mathbf{x}; \mathbf{w}) = \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, \mathbf{y}) - \log Z(\mathbf{x}; \mathbf{w})$$

# CRFs are Factored Log-Linear Models

- For tractability,  $\mathbf{f}(\mathbf{x}, \mathbf{y})$  needs to decompose.  
For **bigram** factorizations:

$$\mathbf{f}(\mathbf{x}_{1:n}, \mathbf{y}_{1:n}) = \sum_{i=1}^n \mathbf{f}(\mathbf{x}, i, y_{i-1}, y_i)$$

- The model form is:

$$\begin{aligned} P(\mathbf{y}|\mathbf{x}_{1:n}; \mathbf{w}) &= \frac{\exp \{ \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, \mathbf{y}) \}}{Z(\mathbf{x}, \mathbf{w})} \\ &= \frac{\exp \{ \sum_{i=1}^n \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, i, y_{i-1}, y_i) \}}{Z(\mathbf{x}, \mathbf{w})} \end{aligned}$$

where

$$Z(\mathbf{x}_{1:n}, \mathbf{w}) = \sum_{\mathbf{z} \in \mathcal{Y}^n} \exp \left\{ \sum_{i=1}^n \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, i, z_{i-1}, z_i) \right\}$$

# Making Predictions with CRFs

- ▶ Given  $\mathbf{w}$ , given  $\mathbf{x}_{1:n}$ , find:

$$\begin{aligned}\operatorname{argmax}_{\mathbf{y} \in \mathcal{Y}^n} P(\mathbf{y} | \mathbf{x}_{1:n}; \mathbf{w}) &= \operatorname{argmax}_{\mathbf{y}} \frac{\exp \left\{ \sum_{i=1}^n \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, i, y_{i-1}, y_i) \right\}}{Z(\mathbf{x}; \mathbf{w})} \\ &= \operatorname{argmax}_{\mathbf{y}} \exp \left\{ \sum_{i=1}^n \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, i, y_{i-1}, y_i) \right\} \\ &= \operatorname{argmax}_{\mathbf{y}} \sum_{i=1}^n \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, i, y_{i-1}, y_i)\end{aligned}$$

- ▶ We can use the Viterbi algorithm

# Making Predictions with CRFs

- ▶ Given  $\mathbf{w}$ , given  $\mathbf{x}_{1:n}$ , find:

$$\begin{aligned}\operatorname{argmax}_{\mathbf{y} \in \mathcal{Y}^n} P(\mathbf{y} | \mathbf{x}_{1:n}; \mathbf{w}) &= \operatorname{argmax}_{\mathbf{y}} \frac{\exp \left\{ \sum_{i=1}^n \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, i, y_{i-1}, y_i) \right\}}{Z(\mathbf{x}; \mathbf{w})} \\ &= \operatorname{argmax}_{\mathbf{y}} \exp \left\{ \sum_{i=1}^n \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, i, y_{i-1}, y_i) \right\} \\ &= \operatorname{argmax}_{\mathbf{y}} \sum_{i=1}^n \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, i, y_{i-1}, y_i)\end{aligned}$$

- ▶ We can use the Viterbi algorithm

# Parameter Estimation in CRFs

- ▶ Given a training set

$$\left\{ (\mathbf{x}^{(1)}, \mathbf{y}^{(1)}), (\mathbf{x}^{(2)}, \mathbf{y}^{(2)}), \dots, (\mathbf{x}^{(m)}, \mathbf{y}^{(m)}) \right\} \quad ,$$

estimate  $\mathbf{w}$

- ▶ Define the conditional log-likelihood of the data:

$$L(\mathbf{w}) = \sum_{k=1}^m \log P(\mathbf{y}^{(k)} | \mathbf{x}^{(k)}; \mathbf{w})$$

- ▶  $L(\mathbf{w})$  measures how well  $\mathbf{w}$  explains the data. A good value for  $\mathbf{w}$  will give a high value for  $P(\mathbf{y}^{(k)} | \mathbf{x}^{(k)}; \mathbf{w})$  for all  $k = 1 \dots m$ .
- ▶ We want  $\mathbf{w}$  that **maximizes**  $L(\mathbf{w})$



# Learning the Parameters of a CRF

- ▶ Recall first lecture on log-linear / maximum-entropy models
- ▶ Find:

$$\mathbf{w}^* = \operatorname{argmax}_{\mathbf{w} \in \mathbb{R}^D} L(\mathbf{w}) - \frac{\lambda}{2} \|\mathbf{w}\|^2$$

where

- ▶ The first term is the log-likelihood of the data
- ▶ The second term is a regularization term, it penalizes solutions with large norm
- ▶  $\lambda$  is a parameter to control the trade-off between fitting the data and model complexity

# Learning the Parameters of a CRF

- Find

$$\mathbf{w}^* = \operatorname{argmax}_{\mathbf{w} \in \mathbb{R}^D} L(\mathbf{w}) - \frac{\lambda}{2} \|\mathbf{w}\|^2$$

- In general there is no analytical solution to this optimization
- We use iterative techniques, i.e. gradient-based optimization
  1. Initialize  $\mathbf{w} = \mathbf{0}$
  2. Take derivatives of  $L(\mathbf{w}) - \frac{\lambda}{2} \|\mathbf{w}\|^2$ , compute gradient
  3. Move  $\mathbf{w}$  in steps proportional to the gradient
  4. Repeat steps 2 and 3 until convergence

## Computing the gradient

$$\begin{aligned}\frac{\partial L(\mathbf{w})}{\partial \mathbf{w}_j} &= \frac{1}{m} \sum_{k=1}^m \mathbf{f}_j(\mathbf{x}^{(k)}, \mathbf{y}^{(k)}) \\ &\quad - \sum_{k=1}^m \sum_{\mathbf{y} \in \mathcal{Y}^*} P(\mathbf{y} | \mathbf{x}^{(k)}; \mathbf{w}) \mathbf{f}_j(\mathbf{x}^{(k)}, \mathbf{y})\end{aligned}$$

where

$$\mathbf{f}(\mathbf{x}, \mathbf{y}) = \sum_{i=1}^n \mathbf{f}_j(\mathbf{x}, i, y_{i-1}, y_i)$$

- ▶ First term: observed mean feature value
- ▶ Second term: expected feature value under current  $\mathbf{w}$

# Computing the gradient

- ▶ The first term is easy to compute, by counting explicitly

$$\frac{1}{m} \sum_{k=1}^m \sum_i \mathbf{f}_j(\mathbf{x}, i, y_{i-1}^{(k)}, y_i^{(k)})$$

- ▶ The second term is more involved,

$$\sum_{k=1}^m \sum_{\mathbf{y} \in \mathcal{Y}^*} P(\mathbf{y} | \mathbf{x}^{(k)}; \mathbf{w}) \sum_i \mathbf{f}_j(\mathbf{x}^{(k)}, i, y_{i-1}, y_i)$$

because it sums over all sequences  $\mathbf{y} \in \mathcal{Y}^*$

# Computing the gradient

- For an example  $(\mathbf{x}^{(k)}, \mathbf{y}^{(k)})$ :

$$\sum_{\mathbf{y} \in \mathcal{Y}^n} P(\mathbf{y} | \mathbf{x}^{(k)}; \mathbf{w}) \sum_{i=1}^n \mathbf{f}_j(\mathbf{x}^{(k)}, i, y_{i-1}, y_i) = \\ \sum_{i=1}^n \sum_{a, b \in \mathcal{Y}} \mu_i^k(a, b) \mathbf{f}_j(\mathbf{x}^{(k)}, i, a, b)$$

where

$$\mu_i^k(a, b) = \sum_{\mathbf{y} \in \mathcal{Y}^n : y_{i-1}=a, y_i=b} P(\mathbf{y} | \mathbf{x}^{(k)}; \mathbf{w})$$

- The quantities  $\mu_i^k$  can be computed efficiently in  $O(nL^2)$  using the forward-backward algorithm

# Forward-Backward for CRFs

- Assume fixed  $\mathbf{x}$ . Calculate in  $O(n|\mathcal{Y}|^2)$

$$\mu_i(a, b) = \sum_{\mathbf{y} \in \mathcal{Y}^n : y_{i-1}=a, y_i=b} P(\mathbf{y} | \mathbf{x}; \mathbf{w}) \quad , \quad 1 \leq i \leq n; \quad a, b \in \mathcal{Y}$$

- **Definition:** forward and backward quantities

$$\alpha_i(a) = \sum_{\mathbf{y}_{1:i} \in \mathcal{Y}^i : y_i=a} \exp \left\{ \sum_{j=1}^i \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, j, y_{j-1}, y_j) \right\}$$

$$\beta_i(b) = \sum_{\mathbf{y}_{i:n} \in \mathcal{Y}^{(n-i+1)} : y_i=b} \exp \left\{ \sum_{j=i+1}^n \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, j, y_{j-1}, y_j) \right\}$$

- $Z = \sum_a \alpha_n(a)$
- $\mu_i(a, b) = \{ \alpha_{i-1}(a) * \exp\{\mathbf{w} \cdot \mathbf{f}(\mathbf{x}, i, a, b)\} * \beta_i(b) * Z^{-1} \}$
- Similarly to Viterbi,  $\alpha_i(a)$  and  $\beta_i(b)$  can be computed efficiently in a recursive manner

# Forward-Backward for CRFs

- Assume fixed  $\mathbf{x}$ . Calculate in  $O(n|\mathcal{Y}|^2)$

$$\mu_i(a, b) = \sum_{\mathbf{y} \in \mathcal{Y}^n: y_{i-1}=a, y_i=b} P(\mathbf{y}|\mathbf{x}; \mathbf{w}) \quad , \quad 1 \leq i \leq n; \quad a, b \in \mathcal{Y}$$

- **Definition:** forward and backward quantities

$$\alpha_i(a) = \sum_{\mathbf{y}_{1:i} \in \mathcal{Y}^i: y_i=a} \exp \left\{ \sum_{j=1}^i \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, j, y_{j-1}, y_j) \right\}$$

$$\beta_i(b) = \sum_{\mathbf{y}_{i:n} \in \mathcal{Y}^{(n-i+1)}: y_i=b} \exp \left\{ \sum_{j=i+1}^n \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, j, y_{j-1}, y_j) \right\}$$

- $Z = \sum_a \alpha_n(a)$
- $\mu_i(a, b) = \{\alpha_{i-1}(a) * \exp\{\mathbf{w} \cdot \mathbf{f}(\mathbf{x}, i, a, b)\} * \beta_i(b) * Z^{-1}\}$
- Similarly to Viterbi,  $\alpha_i(a)$  and  $\beta_i(b)$  can be computed efficiently in a recursive manner

## CRFs: summary so far

- ▶ Log-linear models for sequence prediction,  $P(\mathbf{y}|\mathbf{x}; \mathbf{w})$
- ▶ Computations factorize on label bigrams
- ▶ Model form:

$$\operatorname{argmax}_{\mathbf{y} \in \mathcal{Y}^*} \sum_i \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, i, y_{i-1}, y_i)$$

- ▶ Prediction: uses Viterbi (from HMMs)
- ▶ Parameter estimation:
  - ▶ Gradient-based methods, in practice L-BFGS
  - ▶ Computation of gradient uses forward-backward (from HMMs)



## CRFs: summary so far

- ▶ Log-linear models for sequence prediction,  $P(\mathbf{y}|\mathbf{x}; \mathbf{w})$
- ▶ Computations factorize on label bigrams
- ▶ Model form:

$$\operatorname{argmax}_{\mathbf{y} \in \mathcal{Y}^*} \sum_i \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, i, y_{i-1}, y_i)$$

- ▶ Prediction: uses Viterbi (from HMMs)
- ▶ Parameter estimation:
  - ▶ Gradient-based methods, in practice L-BFGS
  - ▶ Computation of gradient uses forward-backward (from HMMs)
- ▶ **Next Question:** HMMs or CRFs?

# HMMs for sequence prediction

- ▶  $\mathbf{x}$  are the observations,  $\mathbf{y}$  are the hidden states
- ▶ HMMs model the joint distribution  $P(\mathbf{x}, \mathbf{y})$
- ▶ Parameters: (assume  $\mathcal{X} = \{1, \dots, k\}$  and  $\mathcal{Y} = \{1, \dots, l\}$ )
  - ▶  $\pi \in \mathbb{R}^l$ ,  $\pi_a = \Pr(y_1 = a)$
  - ▶  $T \in \mathbb{R}^{l \times l}$ ,  $T_{a,b} = \Pr(y_i = b | y_{i-1} = a)$
  - ▶  $O \in \mathbb{R}^{l \times k}$ ,  $O_{a,c} = \Pr(x_i = c | y_i = a)$
- ▶ Model form

$$P(\mathbf{x}, \mathbf{y}) = \pi_{y_1} O_{y_1, x_1} \prod_{i=2}^n T_{y_{i-1}, y_i} O_{y_i, x_i}$$

- ▶ Parameter Estimation: maximum likelihood by counting events and normalizing

# HMMs and CRFs

► In CRFs:  $\hat{\mathbf{y}} = \text{amax}_{\mathbf{y}} \sum_i \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, i, y_{i-1}, y_i)$

► In HMMs:

$$\begin{aligned}\hat{\mathbf{y}} &= \text{amax}_{\mathbf{y}} \pi_{y_1} O_{y_1, x_1} \prod_{i=2}^n T_{y_{i-1}, y_i} O_{y_i, x_i} \\ &= \text{amax}_{\mathbf{y}} \log(\pi_{y_1} O_{y_1, x_1}) + \sum_{i=2}^n \log(T_{y_{i-1}, y_i} O_{y_i, x_i})\end{aligned}$$

► An HMM can be “ported” into a CRF by setting:

$\mathbf{f}_j(\mathbf{x}, i, y, y')$	$\mathbf{w}_j$
$i = 1 \ \& \ y' = a$	$\log(\pi_a)$
$i > 1 \ \& \ y = a \ \& \ y' = b$	$\log(T_{a,b})$
$y' = a \ \& \ x_i = c$	$\log(O_{a,b})$

► Hence, HMM parameters  $\subset$  CRF parameters

# HMMs and CRFs: main differences

- ▶ Representation:
  - ▶ HMM “features” are tied to the generative process.
  - ▶ CRF features are **very** flexible. They can look at the whole input  $x$  paired with a label bigram  $(y, y')$ .
  - ▶ In practice, for prediction tasks, “good” discriminative features can improve accuracy **a lot**.
- ▶ Parameter estimation:
  - ▶ HMMs focus on explaining the data, both  $x$  and  $y$ .
  - ▶ CRFs focus on the mapping from  $x$  to  $y$ .
  - ▶ A priori, it is hard to say which paradigm is better.
  - ▶ Same dilemma as Naive Bayes vs. Maximum Entropy.

# Structured Prediction

Perceptron, SVMs, CRFs

# Learning Structured Predictors

- ▶ Goal: given training data

$$\{(\mathbf{x}^{(1)}, \mathbf{y}^{(1)}), (\mathbf{x}^{(2)}, \mathbf{y}^{(2)}), \dots, (\mathbf{x}^{(m)}, \mathbf{y}^{(m)})\}$$

learn a predictor  $\mathbf{x} \rightarrow \mathbf{y}$  with small error on unseen inputs

- ▶ In a CRF:

$$\begin{aligned} \operatorname{argmax}_{\mathbf{y} \in \mathcal{Y}^*} P(\mathbf{y} | \mathbf{x}; \mathbf{w}) &= \frac{\exp \left\{ \sum_{i=1}^n \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, i, y_{i-1}, y_i) \right\}}{Z(\mathbf{x}; \mathbf{w})} \\ &= \sum_{i=1}^n \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, i, y_{i-1}, y_i) \end{aligned}$$

- ▶ To predict new values,  $Z(\mathbf{x}; \mathbf{w})$  is not relevant
- ▶ Parameter estimation:  $\mathbf{w}$  is set to maximize likelihood

- ▶ Can we learn  $\mathbf{w}$  more directly, focusing on errors?

# Learning Structured Predictors

- ▶ Goal: given training data

$$\{(\mathbf{x}^{(1)}, \mathbf{y}^{(1)}), (\mathbf{x}^{(2)}, \mathbf{y}^{(2)}), \dots, (\mathbf{x}^{(m)}, \mathbf{y}^{(m)})\}$$

learn a predictor  $\mathbf{x} \rightarrow \mathbf{y}$  with small error on unseen inputs

- ▶ In a CRF:

$$\begin{aligned} \operatorname{argmax}_{\mathbf{y} \in \mathcal{Y}^*} P(\mathbf{y} | \mathbf{x}; \mathbf{w}) &= \frac{\exp \left\{ \sum_{i=1}^n \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, i, y_{i-1}, y_i) \right\}}{Z(\mathbf{x}; \mathbf{w})} \\ &= \sum_{i=1}^n \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, i, y_{i-1}, y_i) \end{aligned}$$

- ▶ To predict new values,  $Z(\mathbf{x}; \mathbf{w})$  is not relevant
  - ▶ Parameter estimation:  $\mathbf{w}$  is set to maximize likelihood
- 
- ▶ Can we **learn**  $\mathbf{w}$  **more directly**, focusing on errors?

# The Structured Perceptron

(Collins, 2002)

- ▶ Set  $\mathbf{w} = \mathbf{0}$
- ▶ For  $t = 1 \dots T$ 
  - ▶ For each training example  $(\mathbf{x}, \mathbf{y})$ 
    1. Compute  $\mathbf{z} = \operatorname{argmax}_{\mathbf{z}} \sum_{i=1}^n \mathbf{f}(\mathbf{x}, i, z_{i-1}, z_i)$
    2. If  $\mathbf{z} \neq \mathbf{y}$

$$\mathbf{w} \leftarrow \mathbf{w} + \sum_i \mathbf{f}(\mathbf{x}, i, y_{i-1}, y_i) - \sum_i \mathbf{f}(\mathbf{x}, i, z_{i-1}, \mathbf{z}_i)$$

- ▶ Return  $\mathbf{w}$



# The Structured Perceptron + Averaging

(Freund and Schapire, 1998) (Collins 2002)

- ▶ Set  $\mathbf{w} = \mathbf{0}$ ,  $\mathbf{w}^a = \mathbf{0}$
- ▶ For  $t = 1 \dots T$ 
  - ▶ For each training example  $(\mathbf{x}, \mathbf{y})$ 
    1. Compute  $\mathbf{z} = \operatorname{argmax}_{\mathbf{z}} \sum_{i=1}^n \mathbf{f}(\mathbf{x}, i, z_{i-1}, z_i)$
    2. If  $\mathbf{z} \neq \mathbf{y}$

$$\mathbf{w} \leftarrow \mathbf{w} + \sum_i \mathbf{f}(\mathbf{x}, i, y_{i-1}, y_i) - \sum_i \mathbf{f}(\mathbf{x}, i, z_{i-1}, \mathbf{z}_i)$$

3.  $\mathbf{w}^a = \mathbf{w}^a + \mathbf{w}$

- ▶ Return  $\mathbf{w}^a / mT$ , where  $m$  is the number of training examples

# Properties of the Perceptron

- ▶ Online algorithm. Often much more efficient than “batch” algorithms
- ▶ If the data is separable, it will converge to parameter values with 0 errors
- ▶ Number of errors before convergence is related to a definition of *margin*. Can also relate margin to generalization properties
- ▶ In practice:
  1. Averaging improves performance **a lot**
  2. Typically reaches a good solution after only a few (say 5) iterations over the training set
  3. Often performs nearly as well as CRFs, or SVMs

## Averaged Perceptron Convergence

Iteration	Accuracy
1	90.79
2	91.20
3	91.32
4	91.47
5	91.58
6	91.78
7	91.76
8	91.82
9	91.88
10	91.91
11	91.92
12	91.96
...	

(results on validation set for a parsing task)

# Margin-based Structured Prediction

- ▶ Let  $\mathbf{f}(\mathbf{x}, \mathbf{y}) = \sum_{i=1}^n \mathbf{f}(\mathbf{x}, i, y_{i-1}, y_i)$
- ▶ Model:  $\operatorname{argmax}_{\mathbf{y} \in \mathcal{Y}^*} \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, \mathbf{y})$
- ▶ Consider an example  $(\mathbf{x}^{(k)}, \mathbf{y}^{(k)})$ :  
 $\exists \mathbf{y} \neq \mathbf{y}^{(k)} : \mathbf{w} \cdot \mathbf{f}(\mathbf{x}^{(k)}, \mathbf{y}^{(k)}) < \mathbf{w} \cdot \mathbf{f}(\mathbf{x}^{(k)}, \mathbf{y}) \implies \text{error}$
- ▶ Let  $\mathbf{y}' = \operatorname{argmax}_{\mathbf{y} \in \mathcal{Y}^* : \mathbf{y} \neq \mathbf{y}^{(k)}} \mathbf{w} \cdot \mathbf{f}(\mathbf{x}^{(k)}, \mathbf{y})$   
Define  $\gamma_k = \mathbf{w} \cdot (\mathbf{f}(\mathbf{x}^{(k)}, \mathbf{y}^{(k)}) - \mathbf{f}(\mathbf{x}^{(k)}, \mathbf{y}'))$
- ▶ The quantity  $\gamma_k$  is a notion of **margin** on example  $k$ :  
 $\gamma_k > 0 \iff$  no mistakes in the example  
high  $\gamma_k \iff$  high confidence

# Margin-based Structured Prediction

- ▶ Let  $\mathbf{f}(\mathbf{x}, \mathbf{y}) = \sum_{i=1}^n \mathbf{f}(\mathbf{x}, i, y_{i-1}, y_i)$
- ▶ Model:  $\operatorname{argmax}_{\mathbf{y} \in \mathcal{Y}^*} \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, \mathbf{y})$
- ▶ Consider an example  $(\mathbf{x}^{(k)}, \mathbf{y}^{(k)})$ :  
 $\exists \mathbf{y} \neq \mathbf{y}^{(k)} : \mathbf{w} \cdot \mathbf{f}(\mathbf{x}^{(k)}, \mathbf{y}^{(k)}) < \mathbf{w} \cdot \mathbf{f}(\mathbf{x}^{(k)}, \mathbf{y}) \implies \text{error}$
- ▶ Let  $\mathbf{y}' = \operatorname{argmax}_{\mathbf{y} \in \mathcal{Y}^* : \mathbf{y} \neq \mathbf{y}^{(k)}} \mathbf{w} \cdot \mathbf{f}(\mathbf{x}^{(k)}, \mathbf{y})$   
Define  $\gamma_k = \mathbf{w} \cdot (\mathbf{f}(\mathbf{x}^{(k)}, \mathbf{y}^{(k)}) - \mathbf{f}(\mathbf{x}^{(k)}, \mathbf{y}'))$
- ▶ The quantity  $\gamma_k$  is a notion of **margin** on example  $k$ :  
 $\gamma_k > 0 \iff$  no mistakes in the example  
high  $\gamma_k \iff$  high confidence

# Margin-based Structured Prediction

- ▶ Let  $\mathbf{f}(\mathbf{x}, \mathbf{y}) = \sum_{i=1}^n \mathbf{f}(\mathbf{x}, i, y_{i-1}, y_i)$
- ▶ Model:  $\operatorname{argmax}_{\mathbf{y} \in \mathcal{Y}^*} \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, \mathbf{y})$
- ▶ Consider an example  $(\mathbf{x}^{(k)}, \mathbf{y}^{(k)})$ :  
 $\exists \mathbf{y} \neq \mathbf{y}^{(k)} : \mathbf{w} \cdot \mathbf{f}(\mathbf{x}^{(k)}, \mathbf{y}^{(k)}) < \mathbf{w} \cdot \mathbf{f}(\mathbf{x}^{(k)}, \mathbf{y}) \implies \text{error}$
- ▶ Let  $\mathbf{y}' = \operatorname{argmax}_{\mathbf{y} \in \mathcal{Y}^* : \mathbf{y} \neq \mathbf{y}^{(k)}} \mathbf{w} \cdot \mathbf{f}(\mathbf{x}^{(k)}, \mathbf{y})$   
Define  $\gamma_k = \mathbf{w} \cdot (\mathbf{f}(\mathbf{x}^{(k)}, \mathbf{y}^{(k)}) - \mathbf{f}(\mathbf{x}^{(k)}, \mathbf{y}'))$
- ▶ The quantity  $\gamma_k$  is a notion of **margin** on example  $k$ :  
 $\gamma_k > 0 \iff$  no mistakes in the example  
high  $\gamma_k \iff$  high confidence

# Mistake-augmented Margins

(Taskar et al, 2004)

$\mathbf{x}^{(k)}$	Jack	London	went	to	Paris
$\mathbf{y}^{(k)}$	PER	PER	-	-	LOC
$\mathbf{y}'$	PER	LOC	-	-	LOC
$\mathbf{y}''$	PER	-	-	-	-
$\mathbf{y}'''$	-	-	PER	PER	-

- ▶ Def:  $e(\mathbf{y}, \mathbf{y}') = \sum_{i=1}^n [y_i \neq y'_i]$   
e.g.,  $e(\mathbf{y}^{(k)}, \mathbf{y}^{(k)}) = 0$ ,  $e(\mathbf{y}^{(k)}, \mathbf{y}') = 1$ ,  $e(\mathbf{y}^{(k)}, \mathbf{y}''') = 5$
- ▶ Def:  $\gamma_{k,\mathbf{y}} = \mathbf{w} \cdot (\mathbf{f}(\mathbf{x}^{(k)}, \mathbf{y}^{(k)}) - \mathbf{f}(\mathbf{x}^{(k)}, \mathbf{y})) - e(\mathbf{y}^{(k)}, \mathbf{y})$
- ▶ Def:  $\gamma_k = \min_{\mathbf{y} \neq \mathbf{y}^{(k)}} \gamma_{k,\mathbf{y}}$

# Structured Hinge Loss

- ▶ Define loss function on example  $k$  as:

$$L(\mathbf{w}, \mathbf{x}^{(k)}, \mathbf{y}^{(k)}) = \max_{\mathbf{y} \in \mathcal{Y}^*} \left( e(\mathbf{y}^{(k)}, \mathbf{y}) - \mathbf{w} \cdot (\mathbf{f}(\mathbf{x}^{(k)}, \mathbf{y}^{(k)}) - \mathbf{f}(\mathbf{x}^{(k)}, \mathbf{y})) \right)$$

- ▶ Leads to an SVM for structured prediction
- ▶ Given a training set, find:

$$\operatorname{argmin}_{\mathbf{w} \in \mathbb{R}^D} \sum_{k=1}^m L(\mathbf{w}, \mathbf{x}^{(k)}, \mathbf{y}^{(k)}) + \frac{\lambda}{2} \|\mathbf{w}\|^2$$



# Regularized Loss Minimization

- ▶ Given a training set  $\{(\mathbf{x}^{(1)}, \mathbf{y}^{(1)}), \dots, (\mathbf{x}^{(m)}, \mathbf{y}^{(m)})\}$ .

Find:

$$\operatorname{argmin}_{\mathbf{w} \in \mathbb{R}^D} \sum_{k=1}^m L(\mathbf{w}, \mathbf{x}^{(k)}, \mathbf{y}^{(k)}) + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

- ▶ Two common loss functions  $L(\mathbf{w}, \mathbf{x}^{(k)}, \mathbf{y}^{(k)})$  :
  - ▶ Log-likelihood loss (CRFs)

$$-\log P(\mathbf{y}^{(k)} \mid \mathbf{x}^{(k)}; \mathbf{w})$$

- ▶ Hinge loss (SVMs)

$$\max_{\mathbf{y} \in \mathcal{Y}^*} \left( e(\mathbf{y}^{(k)}, \mathbf{y}) - \mathbf{w} \cdot (\mathbf{f}(\mathbf{x}^{(k)}, \mathbf{y}^{(k)}) - \mathbf{f}(\mathbf{x}^{(k)}, \mathbf{y})) \right)$$

# Learning Structure Predictors: summary so far

- ▶ Linear models for sequence prediction

$$\operatorname{argmax}_{\mathbf{y} \in \mathcal{Y}^*} \sum_i \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, i, y_{i-1}, y_i)$$

- ▶ Computations factorize on label bigrams
  - ▶ Decoding: using Viterbi
  - ▶ Marginals: using forward-backward
- ▶ Parameter estimation:
  - ▶ Perceptron, Log-likelihood, SVMs
  - ▶ Extensions from classification to the structured case
  - ▶ Optimization methods:
    - ▶ Stochastic (sub)gradient methods (LeCun et al 98) (Shalev-Shwartz et al. 07)
    - ▶ Exponentiated Gradient (Collins et al 08)
    - ▶ SVM Struct (Tsochantaridis et al. 04)
    - ▶ Structured MIRA (McDonald et al 05)

## Generic Structure Prediction

# Sequence Prediction, Beyond Bigrams

- ▶ It is easy to extend the scope of features to  $k$ -grams


$$\mathbf{f}(\mathbf{x}, i, y_{i-k+1:i-1}, y_i)$$

- ▶ In general, think of state  $\sigma_i$  remembering relevant history
  - ▶  $\sigma_i = y_{i-1}$  for bigrams
  - ▶  $\sigma_i = y_{i-k+1:i-1}$  for  $k$ -grams
  - ▶  $\sigma_i$  can be the state at time  $i$  of a deterministic automaton generating  $\mathbf{y}$
- ▶ The structured predictor is

$$\operatorname{argmax}_{\mathbf{y} \in \mathcal{Y}^*} \sum_i \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, i, \sigma_i, y_i)$$

- ▶ Viterbi and forward-backward extend naturally, in  $O(nL^k)$


# Dependency Structures



- ▶ Directed arcs represent **dependencies** between a **head word** and a **modifier word**.
- ▶ E.g.:
  - movie modifies saw,*
  - John modifies saw,*
  - today modifies saw*

# Dependency Parsing: arc-factored models

(McDonald et al. 2005)



- Parse trees decompose into single dependencies  $\langle h, m \rangle$

$$\operatorname{argmax}_{\mathbf{y} \in \mathcal{Y}(\mathbf{x})} \sum_{\langle h, m \rangle \in \mathbf{y}} \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, h, m)$$

- Some features:  $\mathbf{f}_1(\mathbf{x}, h, m) = [ \text{"saw"} \rightarrow \text{"movie"} ]$   
 $\mathbf{f}_2(\mathbf{x}, h, m) = [ \text{distance} = +2 ]$
- Tractable inference algorithms exist (tomorrow's lecture)

# Linear Structured Prediction

- ▶ Sequence prediction (bigram factorization)

$$\operatorname{argmax}_{\mathbf{y} \in \mathcal{Y}(\mathbf{x})} \sum_i \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, i, \mathbf{y}_{i-1}, \mathbf{y}_i)$$

- ▶ Dependency parsing (arc-factored)

$$\operatorname{argmax}_{\mathbf{y} \in \mathcal{Y}(\mathbf{x})} \sum_{\langle h, m \rangle \in y} \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, h, m)$$

- ▶ In general, we can enumerate parts  $r \in \mathbf{y}$

$$\operatorname{argmax}_{\mathbf{y} \in \mathcal{Y}(\mathbf{x})} \sum_{r \in \mathbf{y}} \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, r)$$

# Linear Structured Prediction Framework

- ▶ Abstract models of structures
  - ▶ Input domain  $\mathcal{X}$ , output domain  $\mathcal{Y}$
  - ▶ A choice of factorization,  $r \in \mathbf{y}$
  - ▶ Features:  $\mathbf{f}(\mathbf{x}, r) \rightarrow \mathbb{R}^d$
- ▶ The linear prediction model, with  $\mathbf{w} \in \mathbb{R}^d$

$$\operatorname{argmax}_{\mathbf{y} \in \mathcal{Y}(\mathbf{x})} \sum_{r \in \mathbf{y}} \mathbf{w} \cdot \mathbf{f}(\mathbf{x}, r)$$

- ▶ Generic algorithms for Perceptron, CRF, SVM
  - ▶ Require tractable inference algorithms